

MT Interim Report: Neural Word Alignment

BAILEY PARKER
Johns Hopkins University
bailey@jhu.edu

VIVIAN TSAI
Johns Hopkins University
viv@jhu.edu

WILLIAM WATSON
Johns Hopkins University
billwatson@jhu.edu

Abstract

1 Introduction

2 Data Procurement and Processing

2.1 Symbol Tokenization

2.2 Number Tokenization

2.3 Proper Noun Tokenization

2.4 Lemmatization Techniques

2.5 POS Tagging

3 Model Components

3.1 Preliminary Notation

We begin by stating several operations frequently used in our discussion.

3.1.1 Hadamard Product

To perform element-wise multiplication of two matrices A and B , of equivalent dimensions, we use the hadamard product, defined as the \circ .

$$(A \circ B)_{ij} = A_{ij} \cdot B_{ij} \quad (1)$$

3.1.2 Sigmoid Function

The sigmoid function σ is applied element wise and defined as follows:

$$\sigma(x) = \frac{1}{1 + e^{-x}} \quad (2)$$

3.1.3 Hyperbolic Tangent Function

The hyperbolic tangent function \tanh is defined as follows:

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} \quad (3)$$

and is applied element-wise.

3.1.4 Variables

We define our source sequence as s , target sequence as t , and refer to the i -th source token as s_i . We refer to the j -th target token as t_j . We refer to a matrix ψ whose rows represent values related to source tokens s_i and whose columns refer to values related to target tokens t_j . Hence, ψ is an $|s| \times |t|$ sized matrix.

3.2 Softmax and Log Softmax

The softmax function transforms a vector of values into a probability distribution. Applying the softmax function to an n -dimensional input tensor rescales it so that the elements of the n -dimensional output tensor lie in the range (0,1) and sum to 1.

$$\sigma(\mathbf{x})_i = \frac{e^{x_i}}{\sum_j e^{x_j}} \quad (4)$$

For a matrix ψ , the rows correspond to source words s_i and the columns correspond to target words t_j . In addition, when we softmax with respect to the targets, i.e. $\sigma_t(\psi)$, the softmax is applied per column. If applied per row, we denote this as $\sigma_s(\psi)$.

$$\sigma_t(\psi)_{ij} = \frac{e^{\psi_{ij}}}{\sum_k e^{\psi_{kj}}} \quad (5)$$

$$\sigma_s(\psi)_{ij} = \frac{e^{\psi_{ij}}}{\sum_k e^{\psi_{ik}}} \quad (6)$$

However, it is sometimes better to work in log-space, and use the log softmax operator for numerical stability.

$$\log \sigma_t(\psi)_{ij} = \log \frac{e^{\psi_{ij}}}{\sum_k e^{\psi_{kj}}} \quad (7)$$

$$\log \sigma_s(\psi)_{ij} = \log \frac{e^{\psi_{ij}}}{\sum_k e^{\psi_{ik}}} \quad (8)$$

3.3 Word Embeddings

Word Embedding layers allow for a simple lookup table that stores embeddings of a fixed dictionary and size. More specifically, these layers are often used to store word embeddings and retrieve them using indices. The input to the embedding layer is a list of indices, and the output is the corresponding word embeddings. This allows words to be represented numerically to a set embedding dimension size, and thus can be passed onto later layers.

Word Embeddings can be formulated as a weight matrix W_e , where the vector representation of word w_i is the i -th row of the matrix, and can be represented as $W_e[w_i]$.

$$W_e = \begin{bmatrix} \leftarrow w_1 \rightarrow \\ \vdots \\ \leftarrow w_i \rightarrow \\ \vdots \\ \leftarrow w_n \rightarrow \end{bmatrix} \quad (9)$$

3.4 Gated Recurrent Units (GRU)

Gated Recurrent Units are a more complex formulation to recurrent layers to process sequences, and improve the vanishing gradient problem found in vanilla RNN layers while using less parameters than a Long Short-Term Memory (LSTM) layer[1]. A GRU makes use of 3 gates: r_t reset gate; z_t update gate; n_t new gate. These 3 gates allow for a blending of the hidden state with new input, and are computed as follows for each input x_t in a sequence \mathbf{x} :

$$\begin{aligned} r_t &= \sigma(W_{ir}x_t + b_{ir} + W_{hr}h_{t-1} + b_{hr}) \\ z_t &= \sigma(W_{iz}x_t + b_{iz} + W_{hz}h_{t-1} + b_{hz}) \\ n_t &= \tanh(W_{in}x_t + b_{in} + r_t \circ (W_{hn}h_{t-1} + b_{hn})) \\ h_t &= (1 - z_t) \circ n_t + z_t \circ h_{t-1} \end{aligned} \quad (10)$$

where x_t is input at time t , h_{t-1} is the hidden state of the previous layer at time $t - 1$ or the initial hidden

state at time 0, h_t is the new hidden state at time t , and σ is the sigmoid function.

3.5 Bidirectional GRU

3.6 Alignment Prior

3.7 Batch Matrix Multiplication (BMM)

4 Model Descriptions

4.1 Dot Aligner

4.2 Bidirectional GRU Aligner

4.3 Extensions

5 Loss Function

5.1 Supervised Loss

5.2 Unsupervised Alignment Loss

6 Ground Truth Alignment Results?

7 Current Status

8 Future Work

References

- [1] K. Cho, B. van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio. Learning phrase representations using rnn encoder-decoder for statistical machine translation. *arXiv:1406.1078*, 2014.